




DEADLOCK

Operating System (CSC1036 & INF1036)



Dr. Hasin A. Ahmed
Assistant Professor
Department of Computer Science
Gauhati University



Deadlock

- Computer systems are full of resources that can only be used by one process at a time.
- Common examples include printers and tape drives
- Having two processes simultaneously writing to the printer leads to nonsense.
- Consequently, all operating systems have the ability to (temporarily) grant a process exclusive access to certain resources, both hardware and software.
- For many applications, a process needs exclusive access to not one resource, but several

Deadlock

- Suppose, for example, two processes each want to record a scanned document on a CD.
- Process A requests permission to use the scanner and is granted.
- Process B is programmed differently and requests the CD recorder first and is also granted.
- Now A asks for the CD recorder, but the request is denied until B releases it.
- Unfortunately, instead of releasing the CD recorder B asks for the scanner.
- At this point both processes are blocked and will remain so forever.
- This situation is called a deadlock.

Deadlock: Definition

- *A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.*

Deadlock

- *Deadlocks can occur in a variety of situations besides requesting dedicated I/O devices.*
- *In a database system, for example, a program may have to lock several records it is using, to avoid race conditions.*
- *If process A locks record R1 and process B locks record R2, and then each process tries to lock the other one's record, we also have a deadlock.*
- *A resource can be a hardware device (e.g., a tape drive) or a piece of information (e.g., a locked record in a database)*
- *A resource is anything that can be used by only a single process at any instant of time.*

Deadlock: necessary conditions

- *Coffman et al. (1971) showed that four conditions must hold for there to be a deadlock. All conditions must be hold.*
- *Mutual exclusion condition: Each resource is either currently assigned to exactly one process or is available.*
- *Hold and wait condition: Processes currently holding resources that were granted earlier can request new resources.*
- *No preemption condition: Resources previously granted cannot be forcibly taken away from a process.*
- *Circular wait condition: There must be a circular chain of two or more processes, each of which is waiting for a resource held by the next member of the chain.*

Resource Allocation Graph

- *The graphs have two kinds of nodes: processes, shown as circles, and resources, shown as squares*
- *An arc from a resource node (square) to a process node (circle) means that the resource has previously been requested by, granted to, and is currently held by that process*
- *An arc from a process to a resource means that the process is currently blocked waiting for that resource.*
- *A cycle in the graph means that there is a deadlock involving the processes and resources in the cycle (assuming that there is one resource of each kind)*


Resource Allocation Graph



(a)



(b)



(c)

Figure 3-9. Resource allocation graphs. (a) Holding a resource. (b) Requesting a resource. (c) Deadlock.

Dealing with Deadlock

- 1) *Just ignore the problem altogether.*
- 2) *Detection and recovery. Let deadlocks occur, detect them, and take action.*
- 3) *Dynamic avoidance by careful resource allocation.*
- 4) *Prevention, by structurally negating one of the four conditions necessary to cause a deadlock.*

The Ostrich Algorithm

- *Stick your head in the sand and pretend there is no problem at all.*
- *(Ostriches can run at 60 km/hour and their kick is powerful enough to kill any lion with visions of a big chicken dinner.)*
- *If deadlocks occur on the average once every five years, but system crashes due to hardware failures, compiler errors, and operating system bugs occur once a week, most engineers would not be willing to pay a large penalty in performance or convenience to eliminate deadlocks.*

Deadlock Prevention

- This technique tries to ensure that at least one of the necessary conditions stated by Coffman et al. is not fulfilled.
- **Addressing Mutual exclusion:** Mutual exclusion condition will be fulfilled for non-sharable resources. A sharable resource can never fulfil the mutual exclusion condition, hence will not be part of dealock.
- But we cannot prevent deadlock by denying mutual exclusion because some resources are intrinsically non-sharable.
- **Addressing Hold and Wait:** If we can prevent processes that hold resources from waiting for more resources, we can eliminate deadlocks.

Deadlock Prevention

- One way to achieve this goal is to require all processes to request all their resources before starting execution.
- If everything is available, the process will be allocated whatever it needs and can run to completion.
- If one or more resources are busy, nothing will be allocated and the process would just wait.
- An immediate problem with this approach is that many processes do not know how many resources they will need until after they have started running.
- Processes will keep resources engaged for a long time, reducing efficiency

Deadlock Prevention

- A slightly different way to break the hold-and-wait condition is to require a process requesting a resource to first temporarily release all the resources it currently holds.
- Then it tries to get everything it needs all at once.
- **Addressing no-preemption:** If a process is holding some resources and requests another resource that cannot be immediately allocated to it (that is, the process must wait), then all resources currently being held are preempted.
- In other words, these resources are implicitly released.
- The preempted resources are added to the list of resources for which the process is waiting.
- The process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.

Deadlock Prevention


- Alternatively, if a process requests some resources, we first check whether they are available.
- If they are, we allocate them.
- If they are not, we check whether they are allocated to some other process that is waiting for additional resources.
- If so, we preempt the desired resources from the waiting process and allocate them to the requesting process.
- If the resources are neither available nor held by a waiting process, the requesting process must wait.
- While it is waiting, some of its resources may be preempted, but only if another process requests them.

Deadlock Prevention

- **Addressing Circular wait:** One way to avoid the circular wait is to provide a global numbering of all the resources
- Now the rule is this: processes can request resources whenever they want to, but all requests must be made in numerical order.
- If we go by the following numbering, a process may request first a scanner and then a tape drive, but it may not request first a plotter and then a scanner.
- 1. Imagesetter
- 2. Scanner
- 3. Plotter
- 4. Tape drive
- 5. CD Rom drive

Deadlock Prevention

- With this rule, the resource allocation graph can never have cycles
- Let us see why this is true for the case of two processes
- We can get a deadlock only if A requests resource j and B requests resource i.
- If $i > j$, then A is not allowed to request j because that is lower than what it already has.
- If $i < j$, then B is not allowed to request i because that is lower than what it already has.
- Either way, deadlock is impossible.



Deadlock Prevention

- With multiple processes, the same logic holds.
- At every instant, one of the assigned resources will be highest.
- The process holding that resource will never ask for a resource already assigned.
- It will either finish, or at worst, request even higher numbered resources, all of which are available.
- Eventually, it will finish and free its resources.
- At this point, some other process will hold the highest resource and can also finish.
- In short, there exists a scenario in which all processes finish, so no deadlock is present.

Deadlock Prevention

- Although numerically ordering the resources eliminates the problem of deadlocks, it may be impossible to find an ordering that satisfies everyone.
- Moreover, a perfectly good and available copy of a resource could be inaccessible with such a rule

Deadlock Avoidance

- **Deadlock can avoided not by imposing arbitrary rules on processes but by carefully analyzing each resource request to see if it could be safely granted.**
- **The question arises: is there an algorithm that can always avoid deadlock by making the right choice all the time?**
- **The answer is a qualified yes, but only if certain information is available in advance.**
- **We will discuss about two deadlock avoidance algorithms that can handle two different situations**

Deadlock Avoidance

- A state is safe if the system can allocate resources to each process (up to its maximum) in some order and still avoid a deadlock.
- More formally! a system is in a safe state only if there exists a safe sequence.
- A safe state is not a deadlocked state.
- Conversely, a deadlocked state is an unsafe state. Not all unsafe states are deadlocks, however.
- An unsafe state may lead to a deadlock.

Deadlock Avoidance

- Let us try to understand the concept of safe state:
- *A state is safe if the system can allocate resources to each process (up to its maximum) in some order and still avoid a deadlock.*

Deadlock Avoidance

- To illustrate, let us consider a system with 12 magnetic tape drives and three processes: P0, P1, and P2 .
- Process P0 requires 10 tape drives, process P1 may need as many as 4 tape drives, and process P2 may need up to 9 tape drives.
- Suppose that, at time t_0 , process P0 is holding 5 tape drives, process P1 is holding 2 tape drives, and process P2 is holding 2 tape drives. (Thus, there are 3 free tape drives.)

	Maximum Needs	Currently Holds
P0	10	5
P1	4	2
P2	9	2

Deadlock Avoidance

- *Is the system safe at t_0*

Deadlock Avoidance

- *The answer is yes, the system is in safe state*
- The sequence $\langle P1, P0, P2 \rangle$ satisfies the safety condition
- Initial tape drives: 3
- After execution of P1: 5 (+2)
- After execution of P0: 10 (+5)
- After execution of P2: 12 (+2)

	Maximum Needs	Currently Holds
P0	10	5
P1	4	2
P2	9	2

Deadlock Avoidance

- *Now suppose t_1 process P2 requests and is allocated one more tape drive*
- At this point, only process P1 can be allocated all its tape drives.
- When it returns them, the system will have only 4 available tape drives, not sufficient for rest
- So it is an **unsafe state**

	Maximum Needs	Currently Holds
P0	10	5
P1	4	2
P2	9	3

Deadlock Avoidance

- Given the concept of a safe state, we can define avoidance algorithms that ensure that the system "will never deadlock."
- The idea is simply to ensure that the system will always remain in a safe state.
- Initially, the system is in a safe state.
- Whenever a process requests a resource that is currently available, the system must decide whether the resource can be allocated immediately or whether the process must wait.
- The request is granted only if the allocation leaves the system in a safe state.



Deadlock Avoidance

- **In this scheme, if a process requests a resource that is currently available, it may still have to wait.**
- **Thus, resource utilization may be lower than it would otherwise be.**
- **We are going to discuss about two algorithms for deadlock avoidance**
- **First one is Resource allocation graph based algorithm which is applied on a situation where all resources have single instance**
- **Second one is Banker's algorithm which is used for a system with multi instance resources**

Resource Allocation Graph based Algorithm

- If we have a resource-allocation system with only one instance of each resource type, then this algorithm can be used.
- It uses a variant of the resource-allocation graph for deadlock avoidance.
- In addition to the request and assignment edges of traditional resource allocation graph, a new type of edge called a claim edge is introduced in this variant.
- A claim edge $P_i \rightarrow R_j$ indicates that process P_i may request resource R_j at some time in the future.
- This edge resembles a request edge in direction but is represented in the graph by a dashed line.
- When process P_i requests resource R_j the claim edge $P_i \rightarrow R_j$ is converted to a request edge.

Resource Allocation Graph based Algorithm




Resource Allocation Graph based Algorithm

- Suppose that process P_i requests resource R_j . The request can be granted only if converting the request edge $P_i \rightarrow R_j$ to an assignment edge $R_j \rightarrow P_i$ does not result in the formation of a cycle in the resource-allocation graph.
- Note that we check for safety by using a cycle-detection algorithm.
- An algorithm for detecting a cycle in this graph requires an order of n^2 operations, where n is the number of processes in the system.
- Presence of a cycle will indicate that the system is in unsafe state, and hence the allocation will be delayed.

Resource Allocation Graph based Algorithm

- Although R_2 is currently free, we cannot allocate it to P_2 , since this action will create a cycle in the graph



Banker's Algorithm

- **The resource-allocation-graph based algorithm is not applicable to a resource allocation system with multiple instances of each resource type.**
- **In such situations an algorithm named Banker's algorithm is used.**
- **The name was chosen because the algorithm could be used in a banking system to ensure that the bank never allocated its available cash in such a way that it could no longer satisfy the needs of all its customers.**

Banker's Algorithm

- When a new process enters the system, it must declare the maximum number of instances of each resource type that it may need.
- This number may not exceed the total number of resources in the system.
- When a user requests a set of resources, the system must determine whether the allocation of these resources will leave the system in a safe state.
- If it will, the resources are allocated; otherwise, the process must wait until some other process releases enough resources.

Banker's Algorithm

- Assuming n number of processes and m number of resources, following data structures are used to encode the resource allocation system
- **Available:** A vector of length m indicates the number of available resources of each type. If $\text{Available}[j]$ equals k , there are k instances of resource type R_j available.
- **Max:** An $n \times m$ matrix defines the maximum demand of each process. If $\text{Max}[i][j]$ equals k , then process P_i may request at most k instances of resource type R_j .

Banker's Algorithm

- **Allocation:** An $n \times m$ matrix defines the number of resources of each type currently allocated to each process. If $\text{Allocation}[i][j]$ equals k , then process P_i is currently allocated k instances of resource type R_j .
- **Need:** An $n \times m$ matrix indicates the remaining resource need of each process. If $\text{Need}[i][j]$ equals k then process P_i may need k more instances of resource type R_j to complete its task. Note that $\text{Need}[i][j]$ equals $\text{Max}[i][j] - \text{Allocation}[i][j]$.

Banker's Safety Algorithm

- 1) Let *Work* and *Finish* be vectors of length *m* and *n*, respectively. Initialize *Work* = *Available* and *Finish*[*i*] = *false* for *i* = 0, 1, ..., *n* - 1.
- 2) Find an *i* such that both
 - *Finish*[*i*] == *false*
 - *Need*_{*i*} ≤ *Work*
 - If no such *i* exists, go to step 4.
- 3) *Work* = *Work* + *Allocation*_{*i*}
- *Finish*[*i*] = *true*
- Go to step 2
- 4) If *Finish*[*i*] == *true* for all *i*, then the system is in a safe state, otherwise it is unsafe.

Banker's Safety Algorithm

- 1) This algorithm may require an order of $m \times n^2$ operations to determine whether a state is safe.**

Banker's Resource-Request Algorithm

- 1) If $\text{Request}_i \leq \text{Need}_i$, go to step 2. Otherwise, raise an error condition, since the process has exceeded its maximum claim.
- 2) If $\text{Request}_i \leq \text{Available}$, go to step 3. Otherwise, P_i must wait, since the resources are not available.
- 3) Have the system pretend to have allocated the requested resources to process P_i by modifying the state as follows
 - $\text{Available} = \text{Available} - \text{Request}_i$;
 - $\text{Allocation}_i = \text{Allocation}_i + \text{Request}_i$;
 - $\text{Need}_i = \text{Need}_i - \text{Request}_i$;
- 4) If the resulting resource-allocation state is safe, the transaction is completed, and process P_i is allocated its resources.
- 5) However, if the new state is unsafe, then P_i must wait for Request_i , and the old resource-allocation state is restored.

Banker's Algorithm

Process	Allocation	Max	Available
	A B C	A B C	A B C
P ₀	0 1 0	7 5 3	3 3 2
P ₁	2 0 0	3 2 2	
P ₂	3 0 2	9 0 2	
P ₃	2 1 1	2 2 2	
P ₄	0 0 2	4 3 3	

Banker's Algorithm

Process	Allocation	Max	Available
	A B C	A B C	A B C
P ₀	0 1 0	7 5 3	3 3 2
P ₁	2 0 0	3 2 2	
P ₂	3 0 2	9 0 2	
P ₃	2 1 1	2 2 2	
P ₄	0 0 2	4 3 3	

Process	Need		
	A	B	C
P ₀	7	4	3
P ₁	1	2	2
P ₂	6	0	0
P ₃	0	1	1
P ₄	4	3	1

Banker's Algorithm

$m=3, n=5$

Step 1 of Safety Algo

Work = Available

Work =

3	3	2
---	---	---

0 1 2 3 4

Finish =

false	false	false	false	false
-------	-------	-------	-------	-------

Process	Need		
	A	B	C
P ₀	7	4	3
P ₁	1	2	2
P ₂	6	0	0
P ₃	0	1	1
P ₄	4	3	1

Banker's Algorithm

For $i = 0$ ✗ Step 2
 Need₀ = 7, 4, 3 7,4,3 3,3,2
 Finish [0] is false and Need₀ > Work
 So P₀ must wait But Need ≤ Work

For $i = 1$ ✓ Step 2
 Need₁ = 1, 2, 2 1,2,2 3,3,2
 Finish [1] is false and Need₁ < Work
 So P₁ must be kept in safe sequence

Step 3
 3, 3, 2 2, 0, 0
 Work = Work + Allocation₁
 Work =

A	B	C
5	3	2

 0 1 2 3 4
 Finish =

false	true	false	false	false
-------	------	-------	-------	-------

Process	Need		
	A	B	C
P ₀	7	4	3
P ₁	1	2	2
P ₂	6	0	0
P ₃	0	1	1
P ₄	4	3	1

Banker's Algorithm


For $i = 2$ ✗ Step 2:
 $Need_2 = 6, 0, 0$
 Finish [2] is false and $Need_2 > Work$
 So P_2 must wait

For $i = 3$ ✓ Step 2:
 $Need_3 = 0, 1, 1$
 Finish [3] = false and $Need_3 < Work$
 So P_3 must be kept in safe sequence

Step 3
 $Work = Work + Allocation_3$
 $Work = \begin{matrix} & 5, 3, 2 & & 2, 1, 1 \\ A & B & C \\ \begin{bmatrix} 7 & 4 & 3 \end{bmatrix} \end{matrix}$
 $Finish = \begin{matrix} & 0 & 1 & 2 & 3 & 4 \\ \begin{bmatrix} false & true & false & true & false \end{bmatrix} \end{matrix}$

Process	Need		
	A	B	C
P_0	7	4	3
P_1	1	2	2
P_2	6	0	0
P_3	0	1	1
P_4	4	3	1


Banker's Algorithm

For $i = 4$ Step 2
 $Need_4 = 4, 3, 1$
 $Finish[4] = \text{false}$ and $Need_4 < Work$ 
 So P_4 must be kept in safe sequence

Step 3
 $Work = Work + Allocation_4$
 $Work = \begin{matrix} 7, 4, 3 & 0, 0, 2 \\ A & B & C \\ \boxed{7} & \boxed{4} & \boxed{5} \end{matrix}$
 $Finish = \begin{matrix} 0 & 1 & 2 & 3 & 4 \\ \boxed{\text{false}} & \boxed{\text{true}} & \boxed{\text{false}} & \boxed{\text{true}} & \boxed{\text{true}} \end{matrix}$

Process	Need		
	A	B	C
P_0	7	4	3
P_1	1	2	2
P_2	6	0	0
P_3	0	1	1
P_4	4	3	1

Banker's Algorithm

For $i = 0$
 $Need_0 = 7, 4, 3$  Step 2
 $7, 4, 3$ $7, 4, 5$
 Finish [0] is false and $Need < Work$
 So P_0 must be kept in safe sequence

Step 3
 $7, 4, 5$ $0, 1, 0$
 $Work = Work + Allocation_0$
 $Work =$

A	B	C
7	5	5

 $Finish =$

0	1	2	3	4
true	true	false	true	true

Process	Need		
	A	B	C
P_0	7	4	3
P_1	1	2	2
P_2	6	0	0
P_3	0	1	1
P_4	4	3	1

Banker's Algorithm

For $i = 2$ Step 2
 $Need_2 = 6, 0, 0$
 Finish [2] is false and $Need_2 < Work$
 So P_2 must be kept in safe sequence

Step 3
 $Work = Work + Allocation_2$
 $Work = \begin{matrix} & 7, 5, 5 & & 3, 0, 2 \\ \begin{matrix} A & B & C \end{matrix} \\ \begin{bmatrix} 10 & 5 & 7 \end{bmatrix} \end{matrix}$
 $Finish = \begin{matrix} & 0 & 1 & 2 & 3 & 4 \end{matrix}$
 $\begin{bmatrix} true & true & true & true & true \end{bmatrix}$

Step 4
 $Finish[i] = true$ for $0 \leq i \leq n$
 Hence the system is in Safe state

The safe sequence is P_1, P_3, P_4, P_0, P_2

Process	Need		
	A	B	C
P_0	7	4	3
P_1	1	2	2
P_2	6	0	0
P_3	0	1	1
P_4	4	3	1



Detection and Recovery

- If a system does not employ either a deadlock-prevention or a deadlock avoidance algorithm then a deadlock situation may occur. In this environment, the system must provide:
- An algorithm that examines the state of the system to determine whether a deadlock has occurred
- An algorithm to recover from the deadlock

Detection and Recovery

- If all resources have only a single instance, then we can define a deadlock detection algorithm that uses a variant of the resource-allocation graph, called a wait-for graph.
- We obtain this graph from the resource-allocation graph by removing the resource nodes and collapsing the appropriate edges.
- More precisely, an edge from P_i to P_j in a wait-for graph implies that process P_i is waiting for process P_j to release a resource that P_i needs.
- An edge $P_i \rightarrow P_j$ exists in a wait-for graph if and only if the corresponding resource allocation graph contains two edges $P_i \rightarrow R$ and $R \rightarrow P_j$ for some resource R

Detection and Recovery



Detection and Recovery

- As before, a deadlock exists in the system if and only if the wait-for graph contains a cycle.
- To detect deadlocks, the system needs to maintain the wait-for graph and periodically invoke an algorithm that searches for a cycle in the graph.
- An algorithm to detect a cycle in a graph requires an order of n^2 operations, where n is the number of vertices in the graph.
- On detection of deadlock processes are killed to recover the system.

Detection and Recovery

- The wait-for graph scheme is not applicable to a resource-allocation system with multiple instances of each resource type.
- We turn now to a deadlock detection algorithm that is applicable to such a system.
- The algorithm employs several time-varying data structures that are similar to those used in the banker's algorithm
 - Available
 - Allocation
 - Request

Detection and Recovery

- 1) Let **Work** and **Finish** be vectors of length m and n , respectively. Initialize **Work=Available**. For $i=0, 1, \dots, n-1$, if $\text{Request}_i \neq 0$, then $\text{Finish}[i]=\text{false}$; otherwise, $\text{Finish}[i]=\text{true}$
- 2) Find an index i such that both
 - a. $\text{Finish}[i] == \text{false}$
 - b. $\text{Request}_i \leq \text{Work}$
- If no such i exists, go to step 4.
- 3) $\text{Work} = \text{Work} + \text{Allocation}_i$
 - $\text{Finish}[i] = \text{true}$
 - Go to step 2.
- 4) If $\text{Finish}[i] = \text{false}$, for some i , $0 \leq i < n$, then the system is in a deadlocked state. Moreover, if $\text{Finish}[i] = \text{false}$, then process P_i is deadlocked.

Detection and Recovery

- This algorithm requires an order of $m \times n^2$ operations to detect whether the system is in a deadlocked state.
- When should we invoke the detection algorithm?
- Of course, if the deadlock-detection algorithm is invoked for every resource request, this will incur a considerable overhead in computation time.
- A less expensive alternative is simply to invoke the algorithm at less frequent intervals for example, once per hour or whenever CPU utilization drops below 40 percent.

The background of the slide features a repeating pattern of a network or molecular structure. It consists of small grey circles (nodes) connected by thin grey lines (edges). Some nodes are highlighted with a darker grey or blue dot in the center. The pattern is dense and covers the entire slide area.

thank you

???